MTH321

Numerical analysis of PDEs

Mathematical formulation and methods for mixed Poisson problems

Kainan Zhao 2362073

Rengerile Na 2359838

Suning Peng 2361786



Xi'an Jiaotong-Liverpool University
School of Mathematics and Physics

September 16, 2026

Contents

1 Introduction **2**

2 Problem formulation **3**
2.1 Domain, boundary data and regularity 3
2.2 Manufactured benchmark 4
2.3 Weak formulation and uniqueness 4

3 Finite difference discretisation **5**
3.1 Grid, unknowns and norms 5
3.2 Interior stencil and first-order closure 6
3.3 Ghost-point closure 6
3.4 Dirichlet elimination and matrix scaling 7

4 Stability and error analysis **8**
4.1 Uniform consistency estimates 8
4.2 Comparison principle and discrete solvability 9
4.3 Uniform maximum-norm stability 10

5 Weighted energy formulation and solver implications **12**
5.1 Weighted symmetry and the energy identity 12
5.2 Discrete coercivity and energy-error bounds 13
5.3 Solver implications of the proved matrix structure 14

6 Algorithm and verification protocol **15**
6.1 Assembly, library interfaces and reproducible outputs 15
6.2 Checks that test the equations and checks that test the solve 15
6.3 Residuals, tolerance sweep and acceptance criteria 16

7 Numerical experiments and interpretation **17**
7.1 Four-mesh protocol and global errors 17
7.2 Boundary error, interior error and finite-grid deviations 18
7.3 Raw matrices, weighted structure and sparsity 20
7.4 Weighted-energy convergence 21
7.5 Algebraic-error separation and weighted CG 22
7.6 Measured cost at fixed accuracy thresholds 23

8 Conclusions and limitations **25**

A Prospective coordination plan **28**

B Mathematical drafting declaration and source provenance **29**
B.1 Analytical claim audits available in this draft 29
B.2 Supplied assignment sources 29

Abstract

We study how the right-edge closure affects finite difference solutions of a Poisson problem with Dirichlet data on three sides and Neumann or Robin data on the fourth. A backward derivative row is compared with a centred ghost-point condition eliminated into a boundary PDE row. A discrete comparison principle and a boundary-adapted barrier give mesh-independent maximum-norm estimates. Although both raw boundary defects are $O(h)$, their stability weights differ, yielding $O(h)$ and $O(h^2)$ global error bounds. A separate comparison proves a matching first-order lower bound for the exponential benchmark. Positive diagonal weights produce symmetric positive definite systems and discrete energy estimates. Four-mesh experiments confirm approximately first- and second-order convergence in the trapezoidal discrete L^2, nodal maximum and closure-dependent energy norms. Independent boundary-row, manufactured-solution and uneliminated-system checks test the assembly; tolerance sweeps and weighted CG cross-check the linear solves. Measured assembly-plus-solve costs are compared at fixed nodal-error thresholds, with unsuccessful targets recorded explicitly. The conclusions distinguish proved bounds, observed rates and implementation-dependent timing.

Keywords: Poisson equation; mixed boundary conditions; finite differences; ghost points; comparison principle; error estimates; reproducibility.

1 Introduction

Steady heat conduction and electrostatic models combine prescribed values, prescribed outward derivatives and exchange conditions on different boundary segments. A numerical method must translate all of these conditions into equations consistent with its interior operator. The present report studies this issue through Direction 2 of the supplied Problem Pack: a reproducible Poisson problem on the unit square with Dirichlet values on three sides and Neumann or Robin data on the open right edge.

The direction prescribes Cartesian finite differences and two boundary closures. The first replaces the boundary PDE row by a backward derivative equation. The second eliminates a conceptual ghost point from a centred derivative equation and retains a PDE row at the boundary. Our question is how this change of boundary formulation affects the solution error, the matrix structure and the cost of achieving a named accuracy target. The two formulations differ in the equations they enforce and in their row scaling, as well as in derivative accuracy.

LeVeque [1] develops the finite difference framework for boundary value problems. Within that framework, consistency must be combined with stability before it implies convergence. For the present operators, both raw boundary rows have an $O(h)$ defect. We derive a maximum-norm stability estimate that weights the ghost boundary defect by an additional factor h, explaining its $O(h^2)$ global upper bound. The analysis uses the actual mixed boundary and its corner convention.

The exact function $u^* = e^{x+y}$ determines the forcing and boundary data independently of the assembled matrix. This permits exact error evaluation and a manufactured-solution verification study in the sense of Salari and Knupp [2]. The benchmark verifies numerical methods; it does not calibrate a model against an experimental device.

The analytical contributions are an explicit comparison principle, inverse positivity, uniform maximum-norm estimates, a sharp first-order maximum-error result for the exponential FO benchmark, and a weighted energy identity with discrete L^2 and energy bounds. The weighting is a standard diagonal symmetrisation; no new general CG algorithm is claimed. The computational study tests these results in all four boundary/closure configurations, separates algebraic and discretisation errors on a representative grid, and reports measured cost at fixed accuracy thresholds.

Section 2 states the model and regularity assumptions. Section 3 derives the discrete systems. Sections 4 and 5 establish stability, convergence and weighted structure. Section 6 specifies the implemented algorithm and verification checks. Section 7 presents the numerical evidence and cost comparison. Section 8 relates the measured results to the proved statements and states their limitations.

2 Problem formulation

2.1 Domain, boundary data and regularity

We now specify the selected mixed-boundary benchmark. All variables are nondimensional. Set

$$\begin{aligned} &\Omega = (0,1)^2, \\ &\Gamma_D = (\{0\} \times [0,1]) \cup ([0,1] \times \{0,1\}), \quad \Gamma_B = \{1\} \times (0,1). \end{aligned} \tag{2.1}$$

The governing equations are

$$\begin{cases} -\Delta u = f & \text{in } \Omega, \\ u = g & \text{on } \Gamma_D, \\ \partial_n u + \beta u = s_\beta & \text{on } \Gamma_B. \end{cases} \tag{2.2}$$

The outward normal on Γ_B is $(1,0)$, so $\partial_n u = u_x$ there. The required variants are $(\beta, s_\beta) = (0, q)$ and $(2, r)$. Here q denotes the outward derivative; a conductive heat flux would additionally contain the conductivity and a minus sign. Every corner belongs to the Dirichlet node set. In the weak formulation, traces are understood almost everywhere, and isolated corner points have zero boundary measure.

Figure 1 fixes the boundary ownership used below.

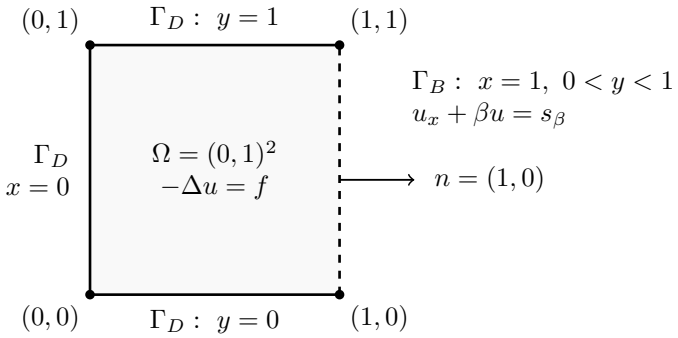


Figure 1: Boundary ownership for the selected benchmark. The solid edges and all four filled corner points carry Dirichlet data. The dashed open right edge carries the derivative/Robin condition; its endpoints are excluded from the boundary-row unknowns. This analytical schematic fixes the outward-normal sign and the corner convention before any computation.

Assumption 2.1 (Smooth solution and consistent data)**.** The coefficient $\beta \geq 0$ is constant and independent of the mesh. The exact solution u extends to a C^6 function on an open neighbourhood of $\overline{\Omega}$. The data satisfy $f = -\Delta u$ on $\overline{\Omega}$, $g = u|_{\Gamma_D}$ and $s_\beta = (u_x + \beta u)|_{\Gamma_B}$. Dirichlet values used in the analysis are imposed exactly. For $\ell = 2, 3, 4$, write

$$M_\ell = \max\{\|\partial_x^\ell u\|_{L^\infty(\overline{\Omega})}, \|\partial_y^\ell u\|_{L^\infty(\overline{\Omega})}\}. \tag{2.3}$$

All asymptotic remainders below are uniform over the relevant grid nodes; their constants may depend on derivative bounds for u, but not on N, h, i or j.

This assumption is verified for the manufactured benchmark. It is an assumption about the solution, rather than a regularity conclusion for arbitrary mixed data; general mixed-boundary problems may have corner singularities.

For a C^1 solution with smooth edge data, a necessary endpoint compatibility is

$$\begin{aligned} \lim_{y\downarrow 0} s_\beta(y) &= \partial_x g(1,0) + \beta g(1,0), \\ \lim_{y\uparrow 1} s_\beta(y) &= \partial_x g(1,1) + \beta g(1,1). \end{aligned} \tag{2.4}$$

Here $\partial_x g$ is the tangential derivative on the bottom or top Dirichlet edge. The identities follow by taking limits of the right-edge condition and differentiating the Dirichlet trace along its horizontal edge. They are necessary conditions for a smooth solution, not extra discrete corner equations. The exponential data below satisfy both.

2.2 Manufactured benchmark

For the prescribed exact solution $u^*(x, y) = e^{x+y}$, direct differentiation gives

$$\begin{aligned} f(x,y) &= -2e^{x+y}, \\ g(0,y) &= e^y, \qquad g(x,0) = e^x, \qquad g(x,1) = e^{x+1}, \\ q(y) &= e^{1+y}, \qquad r(y) = 3e^{1+y}, \quad s_\beta(y) = (1+\beta)e^{1+y}. \end{aligned} \tag{2.5}$$

The corner values at $(0,0), (0,1), (1,0), (1,1)$ are $1, e, e, e^2$, respectively, and $M_2 = M_3 = M_4 = e^2$. The forcing is derived from the continuous solution independently of the discrete matrix. Consequently, substituting sampled exact values into the discrete equations retains the defects that the verification study must detect. Manufactured solutions support code verification through exact error evaluation and systematic refinement [2]; this benchmark does not provide experimental validation of a heat-transfer model.

2.3 Weak formulation and uniqueness

Let Tr denote the Sobolev trace and set

$$V_0 = \{v \in H^1(\Omega) : \operatorname{Tr} v = 0 \text{ on } \Gamma_D\}, \qquad V_g = \{v \in H^1(\Omega) : \operatorname{Tr} v = g \text{ on } \Gamma_D\}.$$

For $u \in V_g$, the weak formulation is

$$a_\beta(u, v) = \ell(v) \qquad (v \in V_0), \tag{2.6}$$

where

$$\begin{aligned} a_\beta(u,v) &= \int_\Omega \nabla u \cdot \nabla v \, dx\, dy + \beta \int_{\Gamma_B} uv \, ds, \\ \ell(v) &= \int_\Omega fv \, dx\, dy + \int_{\Gamma_B} s_\beta v \, ds. \end{aligned} \tag{2.7}$$

The boundary values in these integrals denote traces. This sign convention agrees with the standard weak treatment of Neumann and Robin data in [3].

Lemma 2.2 (Coercivity induced by the left Dirichlet edge)**.** *For every $v \in V_0$,*

$$\|v\|^2_{L^2(\Omega)} \le \tfrac{1}{2}\|v_x\|^2_{L^2(\Omega)}, \qquad \|\operatorname{Tr} v\|^2_{L^2(\Gamma_B)} \le \|v_x\|^2_{L^2(\Omega)}. \tag{2.8}$$

Consequently, for $\beta \ge 0$,

$$a_\beta(v,v) \ge \tfrac{2}{3}\|v\|^2_{H^1(\Omega)}, \qquad \|v\|^2_{H^1(\Omega)} := \|v\|^2_{L^2(\Omega)} + \|\nabla v\|^2_{L^2(\Omega)}. \tag{2.9}$$

Proof. For almost every y, the one-dimensional H^1 slice has an absolutely continuous representative and its left trace is zero. Hence $v(x,y) = \int_0^x v_x(t,y)\,dt$. Cauchy–Schwarz gives

$$|v(x,y)|^2 \le x \int_0^x |v_x(t,y)|^2\,dt.$$

Integrate first in x, interchange the nonnegative integrals, and then integrate in y:

$$\begin{aligned}\|v\|^2_{L^2(\Omega)} &\le \int_0^1 \int_0^1 \left(\int_t^1 x\,dx\right) |v_x(t,y)|^2\,dt\,dy \\ &= \int_0^1 \int_0^1 \tfrac{1-t^2}{2}|v_x(t,y)|^2\,dt\,dy \le \tfrac{1}{2}\|v_x\|^2_{L^2(\Omega)}.\end{aligned}$$

At $x = 1$, the same representation yields $|v(1,y)|^2 \le \int_0^1 |v_x(t,y)|^2\,dt$, proving the trace bound. Finally $\|v\|^2_{H^1} \le \frac{3}{2}\|\nabla v\|^2_{L^2}$, while $a_\beta(v,v) \ge \|\nabla v\|^2_{L^2}$ because $\beta \ge 0$. Combining these inequalities gives (2.9). □

Proposition 2.3 (Solvability for the stated smooth data)**.** *Under Assumption 2.1, the function u is the unique weak solution of* (2.6)*. In particular, both required benchmark variants have the unique solution $u^* = e^{x+y}$.*

Proof. The assumed trace of u shows $u \in V_g$. Green's identity, applicable to this smooth u and $v \in H^1(\Omega)$, gives

$$\int_\Omega \nabla u \cdot \nabla v = \int_\Omega fv + \int_{\partial\Omega} \partial_n u\, v = \int_\Omega fv + \int_{\Gamma_B} (s_\beta - \beta u)v \qquad (v \in V_0).$$

The contribution from Γ_D vanishes because the trace of v is zero there. Moving the Robin term to the left proves (2.6) and establishes existence.

For uniqueness, let $u_1, u_2 \in V_g$ solve (2.6) and put $w = u_1 - u_2 \in V_0$. Subtraction gives $a_\beta(w,v) = 0$ for all $v \in V_0$. Taking $v = w$ yields

$$\|\nabla w\|^2_{L^2(\Omega)} + \beta\|\operatorname{Tr} w\|^2_{L^2(\Gamma_B)} = 0.$$

Lemma 2.2 gives $\frac{2}{3}\|w\|^2_{H^1(\Omega)} \le a_\beta(w,w) = 0$. Therefore $w = 0$ in $H^1(\Omega)$, and $u_1 = u_2$. □

The Dirichlet boundary removes the constant nullspace even when $\beta = 0$. A pure-Neumann compatibility condition and a pressure gauge are therefore not needed for this problem. No existence claim for arbitrary nonsmooth data is used in the subsequent analysis.

3 Finite difference discretisation

3.1 Grid, unknowns and norms

For $N \ge 2$, let $h = 1/N$, $x_i = ih$ and $y_j = jh$. Define the full index set $\mathcal{X}_h = \{0,\dots,N\}^2$ and the disjoint node sets

$$\begin{aligned}\mathcal{I}_h &= \{(i,j) : 1 \le i,j \le N-1\}, \\ \mathcal{B}_h &= \{(N,j) : 1 \le j \le N-1\}, \\ \mathcal{U}_h &= \mathcal{I}_h \cup \mathcal{B}_h, \qquad\qquad \mathcal{D}_h = \mathcal{X}_h \setminus \mathcal{U}_h.\end{aligned} \tag{3.1}$$

A grid function is a real array on $\mathcal{X}_h$. A vector on $\mathcal{U}_h$ will be extended by zero on $\mathcal{D}_h$ whenever used in an error estimate or quadratic form. The ordering is

$$k(i,j) = (i-1) + N(j-1), \qquad 0 \le k < N(N-1) = n_{\mathrm{dof}}. \tag{3.2}$$

There are $(N-1)^2$ interior and $N-1$ right-edge unknowns. Ghost values are not degrees of freedom.

For a full grid function V, write

$$\begin{aligned} \|V\|_{\infty,h} &= \max_{(i,j)\in\mathcal{X}_h} |V_{i,j}|, \\ \|V\|_{2,h}^2 &= h^2 \sum_{i=0}^{N}\sum_{j=0}^{N} \omega_i\omega_j V_{i,j}^2, \qquad \omega_0 = \omega_N = \tfrac{1}{2}, \quad \omega_i = 1 \ (1 \le i \le N-1). \end{aligned} \tag{3.3}$$

For arrays on the interior and open right edge, respectively, also define

$$\|F\|_{I,h}^2 = h^2 \sum_{(i,j)\in\mathcal{I}_h} F_{i,j}^2, \quad \|S\|_{\Gamma,h}^2 = h \sum_{j=1}^{N-1} S_j^2. \tag{3.4}$$

Subscripts ∞, I and ∞, B will denote maxima over $\mathcal{I}_h$ and $\mathcal{B}_h$.

3.2 Interior stencil and first-order closure

The centred second differences are

$$(D_{xx}^h V)_{i,j} = \frac{V_{i+1,j} - 2V_{i,j} + V_{i-1,j}}{h^2}, \qquad (D_{yy}^h V)_{i,j} = \frac{V_{i,j+1} - 2V_{i,j} + V_{i,j-1}}{h^2}.$$

The usual Taylor derivation of these formulas is recalled and quantified in Lemma 4.1; see also [4]. The interior operator is $L_h = -D_{xx}^h - D_{yy}^h$, giving

$$(L_h U)_{i,j} = \frac{4U_{i,j} - U_{i-1,j} - U_{i+1,j} - U_{i,j-1} - U_{i,j+1}}{h^2} = f_{i,j}, \qquad (i,j) \in \mathcal{I}_h. \tag{3.5}$$

Its neighbour at $i = N$ remains an unknown.

The first-order (FO) closure replaces the boundary PDE row with

$$(B_h^\beta U)_j = \left(\frac{1}{h} + \beta\right) U_{N,j} - \frac{1}{h} U_{N-1,j} = s_{\beta,j}, \qquad 1 \le j \le N-1. \tag{3.6}$$

This equation has no tangential coupling. Its diagonal is $1/h$ for Neumann data and $1/h + 2$ for the prescribed Robin data.

3.3 Ghost-point closure

Define the eliminated ghost (G) operator by

$$(G_h^\beta V)_j = \frac{4V_{N,j} - 2V_{N-1,j} - V_{N,j-1} - V_{N,j+1}}{h^2} + \frac{2\beta}{h} V_{N,j}. \tag{3.7}$$

Proposition 3.1 (Exact ghost elimination)**.** *For a fixed boundary index j, the centred boundary condition*

$$\frac{U_{N+1,j} - U_{N-1,j}}{2h} + \beta U_{N,j} = s_{\beta,j} \tag{3.8}$$

and the boundary five-point PDE row

$$4U_{N,j} - U_{N-1,j} - U_{N+1,j} - U_{N,j-1} - U_{N,j+1} = h^2 f_{N,j} \tag{3.9}$$

are equivalent to the single retained equation

$$(G_h^\beta U)_j = f_{N,j} + \frac{2s_{\beta,j}}{h} \tag{3.10}$$

together with the uniquely determined auxiliary value

$$U_{N+1,j} = U_{N-1,j} + 2hs_{\beta,j} - 2h\beta U_{N,j}. \tag{3.11}$$

Proof. Since $h > 0$, multiplying (3.8) by $2h$ and solving for its ghost value gives (3.11). Substitution into (3.9) gives

$$(4 + 2h\beta)U_{N,j} - 2U_{N-1,j} - U_{N,j-1} - U_{N,j+1} = h^2 f_{N,j} + 2hs_{\beta,j}.$$

Dividing by h^2 proves (3.10). Conversely, given retained values satisfying (3.10), define the ghost value by (3.11). The centred boundary condition then holds identically, and multiplication of (3.10) by h^2 and reverse substitution recovers (3.9). This proves both directions and uniqueness of the auxiliary value. □

Thus the ghost method imposes a boundary PDE row after elimination, whereas the first method imposes a derivative row. Both have exactly $N(N-1)$ retained equations and unknowns.

3.4 Dirichlet elimination and matrix scaling

For $c \in \{\mathrm{FO}, \mathrm{G}\}$, let A_h^c denote the reduced matrix obtained from (3.5) and the corresponding raw boundary row (3.6) or (3.10). If $\mathcal{A}_h^c$ is the stencil matrix with columns for all nodes, the reduced system is

$$A_h^c \boldsymbol{U}_h^c = \boldsymbol{b}_h^c, \qquad A_h^c = (\mathcal{A}_h^c)_{\mathcal{U}_h\mathcal{U}_h}, \qquad \boldsymbol{b}_h^c = \boldsymbol{F}_h^c - (\mathcal{A}_h^c)_{\mathcal{U}_h\mathcal{D}_h}\boldsymbol{g}_h. \tag{3.12}$$

Here $\boldsymbol{F}_h^c$ is f in the interior and is s_β (FO) or $f + 2s_\beta/h$ (G) on the right edge. In particular, a known entry $a_{kD}g_D$ contributes $-a_{kD}g_D$ to the right-hand side. The ghost rows at $j = 1$ and $j = N-1$ receive $g(1,0)/h^2$ and $g(1,1)/h^2$, respectively. For $N = 2$ both contributions enter the same row. Neither corner is assigned a derivative row.

Across the interface between $\mathcal{I}_h$ and $\mathcal{B}_h$, the interior coefficient is $-1/h^2$, while the reverse coefficient is $-1/h$ (FO) or $-2/h^2$ (G). Hence both raw matrices are nonsymmetric for $N \geq 2$. Their symmetrised forms will be established in Proposition 5.1, rather than inferred from the continuous Laplacian.

Proposition 3.2 (Analytical sparsity counts)**.** *For $N \geq 2$ and $\beta \geq 0$, the numbers of nonzero entries in the reduced matrices are*

$$\begin{aligned} \mathrm{nnz}(A_h^{\mathrm{FO}}) &= (N-1)(5N-6), \\ \mathrm{nnz}(A_h^{\mathrm{G}}) &= (N-1)(5N-6) + 2(N-2). \end{aligned} \tag{3.13}$$

Both counts grow linearly with $n_{\mathrm{dof}} = N(N-1)$.

Proof. Every interior row contains a diagonal and an east unknown, contributing $2(N-1)^2$ entries. West unknowns occur for $i \geq 2$, contributing $(N-2)(N-1)$ entries. Each of the two vertical directions contributes $(N-2)(N-1)$ entries. Thus the interior rows contain $2(N-1)^2 + 3(N-2)(N-1) = (N-1)(5N-8)$ entries. Every boundary row contains its diagonal and west unknown, adding $2(N-1)$ entries. The FO closure has no further entries, giving its count in (3.13). The G closure also has $N-2$ unknown neighbours in each vertical direction, giving the additional $2(N-2)$ entries. All counted off-diagonals are strictly negative and all diagonals are strictly positive; none vanish for the stated β. Dirichlet entries have already

been removed. These counts include $N = 2$, when the additional vertical-boundary count is zero. Every unknown contributes a nonzero diagonal, giving $\mathrm{nnz}(A_h^c) \geq n_{\mathrm{dof}}$. The formulas also give $\mathrm{nnz}(A_h^c) \leq 5n_{\mathrm{dof}}$ for both closures and all $N \geq 2$. These two bounds prove the stated linear growth. □

These are counts of the mathematically specified operators, not a spy plot of a computed matrix. They do not determine factorisation fill or wall-clock time.

4 Stability and error analysis

4.1 Uniform consistency estimates

Let u_h denote u sampled on the full grid and let $\boldsymbol{u}_h$ be its restriction to the unknowns. Define

$$\boldsymbol{\tau}_h^c = A_h^c \boldsymbol{u}_h - \boldsymbol{b}_h^c. \tag{4.1}$$

Exact Dirichlet elimination means that these entries are precisely

$$\begin{aligned} \tau_{I,i,j} &= (L_h u_h)_{i,j} - f_{i,j}, \\ \tau_{B,j}^{\mathrm{FO}} &= (B_h^\beta u_h)_j - s_{\beta,j}, \\ \tau_{B,j}^{\mathrm{G}} &= (G_h^\beta u_h)_j - f_{N,j} - 2s_{\beta,j}/h. \end{aligned} \tag{4.2}$$

This definition fixes the row scaling in every subsequent estimate.

Lemma 4.1 (Local defects and uniform bounds)**.** *Under Assumption 2.1, the following expansions hold, with derivatives evaluated at the relevant node:*

$$\tau_I = -\frac{h^2}{12}(u_{xxxx} + u_{yyyy}) + O(h^4), \tag{4.3}$$

$$\tau_B^{\mathrm{FO}} = -\frac{h}{2}u_{xx} + \frac{h^2}{6}u_{xxx} + O(h^3), \tag{4.4}$$

$$\tau_B^{\mathrm{G}} = \frac{h}{3}u_{xxx} - \frac{h^2}{12}(u_{xxxx} + u_{yyyy}) + O(h^3). \tag{4.5}$$

For every $N \geq 2$, the explicit bounds are

$$\begin{aligned} \|\tau_I\|_{\infty,I} &\leq \frac{M_4}{6}h^2, \\ \|\tau_B^{\mathrm{FO}}\|_{\infty,B} &\leq \frac{M_2}{2}h, \\ \|\tau_B^{\mathrm{G}}\|_{\infty,B} &\leq \frac{M_3}{3}h + \frac{M_4}{12}h^2. \end{aligned} \tag{4.6}$$

Proof. Taylor's theorem along the x grid line gives, at an interior node,

$$u(x_i \pm h, y_j) = \sum_{\ell=0}^{5} \frac{(\pm h)^\ell}{\ell!}\partial_x^\ell u(x_i, y_j) + R_\pm, \qquad |R_\pm| \leq C_6 h^6/720,$$

where C_6 is a uniform bound for the pure sixth derivatives on $\overline{\Omega}$. Adding the expansions cancels all odd terms and yields $D_{xx}^h u = u_{xx} + h^2 u_{xxxx}/12 + O(h^4)$. The same computation in y, followed by $f = -\Delta u$, proves (4.3). Using Taylor's theorem only through order three instead bounds each centred second-derivative defect by $M_4 h^2/12$, proving the first inequality in (4.6).

At a right-edge node, write $u = u(1, y_j)$ and $u_W = u(1-h, y_j)$. Taylor's theorem gives

$$u_W = u - hu_x + \frac{h^2}{2}u_{xx} - \frac{h^3}{6}u_{xxx} + O(h^4).$$

Substitution in $(u - u_W)/h + \beta u - (u_x + \beta u)$ proves (4.4). More directly, the second-order Taylor remainder has magnitude at most $M_2 h^2/2$, giving $|\tau_B^{\mathrm{FO}}| \le M_2 h/2$.

For the ghost row, expand the inward and tangential terms separately:

$$2u - 2u_W = 2hu_x - h^2 u_{xx} + \frac{h^3}{3}u_{xxx} - \frac{h^4}{12}u_{xxxx} + O(h^5),$$
$$2u - u(1, y_j - h) - u(1, y_j + h) = -h^2 u_{yy} - \frac{h^4}{12}u_{yyyy} + O(h^6).$$

After division by h^2 and addition of $2\beta u/h$, this becomes

$$G_h^\beta u_h = \frac{2}{h}(u_x + \beta u) - \Delta u + \frac{h}{3}u_{xxx} - \frac{h^2}{12}(u_{xxxx} + u_{yyyy}) + O(h^3).$$

The first two terms equal $2s_\beta/h + f$, proving (4.5). For its uniform bound, use the exact decomposition

$$\tau_B^{\mathrm{G}} = \left[\frac{2}{h^2}(u - u_W - hu_x) + u_{xx}\right] + (u_{yy} - D_{yy}^h u_h).$$

The third-order Taylor remainder for u_W bounds the bracket by $M_3 h/3$. The centred tangential defect is at most $M_4 h^2/12$. The triangle inequality proves the final bound. Every Taylor segment used here lies in $\overline{\Omega}$, including when $N = 2$; bounded derivatives therefore make all constants independent of the node and mesh. □

Corollary 4.2 (Exact benchmark identities)**.** *For $u^* = e^{x+y}$, the defects satisfy*

$$\begin{aligned}
\frac{\tau_I}{u^*} &= 2 + \frac{4 - 2e^h - 2e^{-h}}{h^2}, \\
\frac{\tau_B^{\mathrm{FO}}}{u^*} &= \frac{1 - e^{-h}}{h} - 1, \\
\frac{\tau_B^{\mathrm{G}}}{u^*} &= \frac{4 - 3e^{-h} - e^h}{h^2} - \frac{2}{h} + 2.
\end{aligned} \tag{4.7}$$

The limits of these ratios divided by h^2, h and h, respectively, are $-1/6$, $-1/2$ and $1/3$ as $h \to 0$.

Proof. Relative to a node's exact value, its east and north neighbours contribute e^h, while its west and south neighbours contribute e^{-h}. Insert these ratios and the data (2.5) into (4.2). The Robin terms cancel exactly. Expanding $e^{\pm h}$ through the first nonzero defect term gives the stated limits. □

Although the centred derivative in (3.8) has an $O(h^2)$ defect, the raw eliminated PDE row has an $O(h)$ defect. The global effect of that row must be established using its boundary stability weight.

4.2 Comparison principle and discrete solvability

Lemma 4.3 (Discrete comparison principle)**.** *Let $N \ge 2$ and $\beta \ge 0$. Suppose a grid function V satisfies $V \ge 0$ on $\mathcal{D}_h$ and $L_h V \ge 0$ on $\mathcal{I}_h$. If either $B_h^\beta V \ge 0$ on $\mathcal{B}_h$ or $G_h^\beta V \ge 0$ there, then $V \ge 0$ on the whole grid.*

Proof. Suppose instead that the minimum of V is $m < 0$. The finite grid contains a node attaining m, which cannot belong to $\mathcal{D}_h$. At an interior minimum $P = (i, j)$,

$$0 \le (L_h V)_P = \frac{1}{h^2} \sum_{Q \in \{(i-1,j),(i+1,j),(i,j-1),(i,j+1)\}} (m - V_Q) \le 0.$$

Each summand is nonpositive and their sum is zero, so every neighbour equals m. In particular, the west neighbour is another minimum. Repeat this argument with decreasing i until reaching $(0, j)$, contradicting the nonnegative Dirichlet value.

If the initial minimum is at (N, j), the FO row gives

$$0 \le (B_h^\beta V)_j = \frac{m - V_{N-1,j}}{h} + \beta m \le 0.$$

For $\beta > 0$ the final inequality is strict, an immediate contradiction. For $\beta = 0$, equality forces $V_{N-1,j} = m$; since $N \ge 2$, this is an interior node and the preceding argument applies. For the ghost row,

$$0 \le (G_h^\beta V)_j = \frac{2(m - V_{N-1,j}) + (m - V_{N,j-1}) + (m - V_{N,j+1})}{h^2} + \frac{2\beta}{h} m \le 0.$$

Again $\beta > 0$ gives a contradiction. If $\beta = 0$, every weighted nonpositive term must vanish; in particular $V_{N-1,j} = m$, so the same interior argument applies. These cases cover every possible minimum. □

We use the standard inverse-positive characterisation of nonsingular M-matrices; see Berman and Plemmons [5]. The properties needed for these particular matrices are proved directly below.

Proposition 4.4 (Invertibility and inverse positivity)**.** *For $N \ge 2$, $\beta \ge 0$ and either closure, A_h^c is invertible, has nonpositive off-diagonal entries, and satisfies $(A_h^c)^{-1} \ge 0$ entrywise. Thus it has the inverse-positive characterisation of a nonsingular M-matrix.*

Proof. The sign of every off-diagonal entry follows directly from the stencils. If $A_h^c \boldsymbol{V} = 0$, extend $\boldsymbol{V}$ by zero on $\mathcal{D}_h$. Its full-grid interior and boundary rows are zero. Lemma 4.3 applied to both V and $-V$ gives $V = 0$. The square matrix A_h^c therefore has a trivial kernel and is invertible. For any vector $\boldsymbol{d} \ge 0$, its solution $\boldsymbol{V} = (A_h^c)^{-1}\boldsymbol{d}$, extended by zero, satisfies the hypotheses of that lemma and hence is nonnegative. Taking $\boldsymbol{d}$ to be each coordinate vector proves that every column of $(A_h^c)^{-1}$ is nonnegative. □

4.3 Uniform maximum-norm stability

Lemma 4.5 (A boundary-adapted barrier)**.** *Let $\psi_{i,j} = \psi(x_i)$, where $\psi(x) = 1 + x - x^2/4$. For $N \ge 2$, $1 \le \psi_{i,j} \le 7/4$, and the full-grid operators satisfy*

$$\begin{aligned} L_h\psi &= \tfrac{1}{2}, \\ B_h^\beta\psi &= \tfrac{1}{2} + \tfrac{h}{4} + \tfrac{7\beta}{4} \ge \tfrac{1}{2}, \\ G_h^\beta\psi &= \tfrac{1}{h} + \tfrac{1}{2} + \tfrac{7\beta}{2h} \ge \tfrac{1}{h}. \end{aligned} \tag{4.8}$$

Proof. Since $\psi'(x) = 1 - x/2 > 0$ on $[0, 1]$, its endpoint values give the stated range. The centred differences are exact on this quadratic and ψ does not depend on y, so $L_h\psi = -\psi'' = 1/2$. Moreover $\psi(1) = 7/4$ and $\psi(1) - \psi(1 - h) = h/2 + h^2/4$. Substitution into the FO row yields the second identity. In the ghost row the tangential differences vanish, leaving $2[\psi(1) - \psi(1-h)]/h^2 + 2\beta\psi(1)/h$, which gives the third identity. The inequalities use $\beta \ge 0$. □

Proposition 4.6 (Uniform stability in the raw row scaling)**.** *Let $N \geq 2$, $\beta \geq 0$ and $V = 0$ on $\mathcal{D}_h$. Suppose $L_h V = F$ on $\mathcal{I}_h$ and the selected boundary operator applied to V equals S on $\mathcal{B}_h$. Set $F_0 = \|F\|_{\infty,I}$ and $S_0 = \|S\|_{\infty,B}$. Then*

$$\|V\|_{\infty,h} \leq \tfrac{7}{2}(F_0 + S_0) \qquad \text{for FO,} \tag{4.9}$$

$$\|V\|_{\infty,h} \leq \tfrac{7}{4}(2F_0 + hS_0) \qquad \text{for G.} \tag{4.10}$$

The constants are independent of N and $\beta \geq 0$.

Proof. Use the full-grid barrier of Lemma 4.5. For FO, set $a = 2(F_0 + S_0)$. On $\mathcal{D}_h$, $a\psi \pm V = a\psi \geq 0$. In the interior, $L_h(a\psi \pm V) = a/2 \pm F \geq a/2 - F_0 \geq 0$. On the right edge, $B_h^\beta(a\psi \pm V) \geq a/2 - S_0 \geq 0$. Lemma 4.3 gives $a\psi - V \geq 0$ and $a\psi + V \geq 0$, so $|V| \leq a\psi \leq 7a/4$, proving (4.9).

For G, choose $a = 2F_0 + hS_0$. The Dirichlet values are again nonnegative, the interior rows satisfy $a/2 - F_0 = hS_0/2 \geq 0$, and the boundary rows satisfy $G_h^\beta(a\psi \pm V) \geq a/h - S_0 = 2F_0/h \geq 0$. The same comparison argument yields $|V| \leq 7a/4$ and proves (4.10). The barrier has its own nonzero Dirichlet values; the proof uses full-grid operators, not a reduced matrix multiplication that discards those values. □

Theorem 4.7 (Maximum-norm convergence)**.** *Under Assumption 2.1, let U_h^c be the unique discrete solution, including its exact Dirichlet values, and set $e_h^c = U_h^c - u_h$. For all $N \geq 2$,*

$$\|e_h^{\mathrm{FO}}\|_{\infty,h} \leq \tfrac{7}{2}\left(\tfrac{M_4}{6}h^2 + \tfrac{M_2}{2}h\right), \tag{4.11}$$

$$\|e_h^{\mathrm{G}}\|_{\infty,h} \leq \tfrac{7}{4}\left(\tfrac{M_4+M_3}{3}h^2 + \tfrac{M_4}{12}h^3\right). \tag{4.12}$$

Consequently, these maximum-norm errors are $O(h)$ and $O(h^2)$, respectively, as $N \to \infty$.

Proof. Proposition 4.4 gives existence and uniqueness of each discrete solution. On $\mathcal{D}_h$, $e_h^c = 0$. Subtraction of (4.1) from the discrete system gives

$$A_h^c \boldsymbol{e}_h^c = -\boldsymbol{\tau}_h^c. \tag{4.13}$$

Thus its full-grid interior equation is $L_h e_h^c = -\tau_I$ and its boundary equation is $B_h^\beta e_h^{\mathrm{FO}} = -\tau_B^{\mathrm{FO}}$ or $G_h^\beta e_h^{\mathrm{G}} = -\tau_B^{\mathrm{G}}$. Apply Proposition 4.6 and then Lemma 4.1 to obtain (4.11) and (4.12). Since $h \leq 1/2$ and M_ℓ is independent of h, the first bound is at most $\frac{7}{2}(M_4/12 + M_2/2)h$, and the second at most $\frac{7}{4}[(M_4 + M_3)/3 + M_4/24]h^2$. This proves uniform convergence with the stated upper orders. □

Theorem 4.8 (Sharp first-order maximum error for the exponential benchmark)**.** *Let $u^* = e^{x+y}$ and fix $\beta \geq 0$. For the exact FO discrete solution and every $N \geq 2$,*

$$\max_{1 \leq j \leq N-1} |e_{N,j}^{\mathrm{FO}}| \geq C_\beta h, \qquad C_\beta := \frac{e \sinh \pi}{2\sqrt{2}\, m_\beta} > 0, \quad m_\beta := \pi \cosh \pi + \beta \sinh \pi. \tag{4.14}$$

Together with Theorem 4.7, this implies $\|e_h^{\mathrm{FO}}\|_{\infty,h} = \Theta(h)$ and $\max_j |e_{N,j}^{\mathrm{FO}}| = \Theta(h)$ for each fixed β, in particular for both required variants.

Proof. For the exponential, the interior identity (4.7) and $\cosh h - 1 \geq h^2/2$ give

$$\frac{\tau_I}{u^*} = 2 - \frac{4(\cosh h - 1)}{h^2} \leq 0.$$

At a boundary node, cancellation of the Robin term and the fundamental theorem of calculus give the exact formula

$$-\tau_{B,j}^{\mathrm{FO}} = \frac{1}{h}\int_{1-h}^{1} (z - 1 + h) u_{xx}^*(z, y_j)\, dz.$$

Since $y_j \geq h$ and $z \geq 1-h$, $u^*_{xx}(z, y_j) = e^{z+y_j} \geq e$. The integral of $z - 1 + h$ is $h^2/2$, so $-\tau^{\mathrm{FO}}_{B,j} \geq eh/2$.

Set $\phi_{i,j} = \sinh(\pi x_i)\sin(\pi y_j)$. It vanishes on $\mathcal{D}_h$ and is nonnegative on the full grid. Direct use of the addition formulas yields

$$(L_h\phi)_{i,j} = \frac{4 - 2\cosh(\pi h) - 2\cos(\pi h)}{h^2}\phi_{i,j} \leq 0.$$

The sign follows from $\cosh t + \cos t = 2\sum_{k=0}^{\infty} t^{4k}/(4k)! \geq 2$ for real t. On the right edge,

$$(B^{\beta}_h\phi)_j = \left[\frac{\sinh\pi - \sinh(\pi(1-h))}{h} + \beta\sinh\pi\right]\sin(\pi y_j) \leq m_\beta.$$

Indeed the backward quotient equals the average of $\pi\cosh(\pi x)$ over $[1-h, 1]$, which is at most $\pi\cosh\pi$, and $0 \leq \sin(\pi y_j) \leq 1$.

Choose $a = e/(2m_\beta) > 0$ and define $V = e^{\mathrm{FO}}_h - ah\phi$. It has zero Dirichlet values. The error equation gives

$$L_hV = -\tau_I - ahL_h\phi \geq 0, \qquad B^{\beta}_hV = -\tau^{\mathrm{FO}}_B - ahB^{\beta}_h\phi \geq eh/2 - ahm_\beta = 0.$$

Lemma 4.3 implies $e^{\mathrm{FO}}_h \geq ah\phi$. Take a boundary index j closest to $N/2$. It belongs to $\{1, \ldots, N-1\}$ and satisfies $|y_j - 1/2| \leq h/2 \leq 1/4$, so $\sin(\pi y_j) \geq 1/\sqrt{2}$. Therefore $e^{\mathrm{FO}}_{N,j} \geq ah\sinh\pi/\sqrt{2} = C_\beta h$. The maximum-norm upper bound in Theorem 4.7 also bounds the boundary maximum by a constant times h. These two-sided bounds prove both $\Theta(h)$ statements. □

Remark 4.9 (Meaning of the convergence claim). The factor h on the ghost boundary defect in (4.10) converts its $O(h)$ local size into an $O(h^2)$ contribution to the solution error. Theorem 4.7 alone establishes upper bounds, not nonzero leading constants. Theorem 4.8 provides a separate lower bound for the exponential FO maximum errors; no such lower bound is proved here for G. For affine solutions the defect bounds vanish for both closures, and for quadratic solutions they vanish for G, illustrating that an upper order does not guarantee a nonzero error for every smooth solution.

5 Weighted energy formulation and solver implications

5.1 Weighted symmetry and the energy identity

For either closure define the positive diagonal matrix W^c_h by

$$(W^c_h)_{kk} = \begin{cases} 1, & k \text{ corresponds to } \mathcal{I}_h, \\ 1/h, & k \text{ corresponds to } \mathcal{B}_h,\ c = \mathrm{FO}, \\ 1/2, & k \text{ corresponds to } \mathcal{B}_h,\ c = \mathrm{G}, \end{cases} \qquad K^c_h = W^c_hA^c_h. \tag{5.1}$$

For a vector V, extended by zero on $\mathcal{D}_h$, write

$$\begin{aligned} S_x(V) &= \sum_{j=1}^{N-1}\sum_{i=0}^{N-1}(V_{i+1,j} - V_{i,j})^2, \\ S_{y,I}(V) &= \sum_{i=1}^{N-1}\sum_{j=0}^{N-1}(V_{i,j+1} - V_{i,j})^2, \\ S_{y,B}(V) &= \sum_{j=0}^{N-1}(V_{N,j+1} - V_{N,j})^2. \end{aligned} \tag{5.2}$$

Proposition 5.1 (Symmetry, energy identity and positive definiteness)**.** *For $N \geq 2$ and $\beta \geq 0$, K_h^c is real symmetric positive definite. For every real vector $\boldsymbol{V}$,*

$$h^2 \boldsymbol{V}^T K_h^c \boldsymbol{V} = S_x(V) + S_{y,I}(V) + \theta_c S_{y,B}(V) + \beta h \sum_{j=1}^{N-1} V_{N,j}^2, \quad \theta_{\mathrm{FO}} = 0, \quad \theta_{\mathrm{G}} = \tfrac{1}{2}. \tag{5.3}$$

Proof. Interior–interior off-diagonal pairs already agree. On each interior–boundary pair, multiplying the boundary coefficient by $1/h$ (FO) or $1/2$ (G) makes both coefficients $-1/h^2$. In G the two tangential boundary coefficients both become $-1/(2h^2)$; in FO both are zero. There are no other off-diagonal pairs, so K_h^c is symmetric.

For a fixed j, abbreviate $v_i = V_{i,j}$, with $v_0 = 0$. Expansion of the squares proves the finite summation identity

$$\sum_{i=1}^{N-1} v_i(2v_i - v_{i-1} - v_{i+1}) + v_N(v_N - v_{N-1}) = \sum_{i=0}^{N-1} (v_{i+1} - v_i)^2. \tag{5.4}$$

Indeed both sides equal $2\sum_{i=1}^{N-1} v_i^2 + v_N^2 - 2\sum_{i=1}^{N-1} v_i v_{i+1}$. For a vertical sequence z_j with $z_0 = z_N = 0$, the same expansion gives

$$\sum_{j=1}^{N-1} z_j(2z_j - z_{j-1} - z_{j+1}) = \sum_{j=0}^{N-1} (z_{j+1} - z_j)^2. \tag{5.5}$$

Multiplication of the scaled matrix by h^2 leaves the usual interior rows and the following boundary rows:

$$\begin{cases} V_{N,j} - V_{N-1,j} + \beta h V_{N,j}, & c = \mathrm{FO}, \\ V_{N,j} - V_{N-1,j} + \frac{1}{2}(2V_{N,j} - V_{N,j-1} - V_{N,j+1}) + \beta h V_{N,j}, & c = \mathrm{G}. \end{cases}$$

Apply (5.4) to every horizontal sequence and (5.5) to each interior column. In G also apply the vertical identity to the right column with factor $1/2$. Summing the remaining diagonal Robin terms proves (5.3), including the boundary edges adjoining the Dirichlet corners.

Every term on the right of (5.3) is nonnegative. If the quadratic form is zero, $S_x(V) = 0$, so every horizontal difference is zero. Starting from $V_{0,j} = 0$ then gives $V_{i,j} = 0$ for all unknowns. The symmetric quadratic form is therefore strictly positive for every nonzero vector, proving positive definiteness. □

It follows that

$$\|V\|_{E,h,c} = \left(h^2 \boldsymbol{V}^T K_h^c \boldsymbol{V}\right)^{1/2} \tag{5.6}$$

is a norm on the unknown vectors with zero Dirichlet extension.

5.2 Discrete coercivity and energy-error bounds

Lemma 5.2 (Discrete Poincaré and right-edge trace bounds)**.** *For $V = 0$ on $\mathcal{D}_h$, $N \geq 2$ and $\beta \geq 0$,*

$$\|V\|_{2,h}^2 \leq S_x(V) \leq \|V\|_{E,h,c}^2, \qquad \|V_B\|_{\Gamma,h}^2 \leq S_x(V) \leq \|V\|_{E,h,c}^2. \tag{5.7}$$

Proof. For fixed j, put $d_{r,j} = V_{r+1,j} - V_{r,j}$. Since $V_{0,j} = 0$, $V_{i,j} = \sum_{r=0}^{i-1} d_{r,j}$. The Cauchy–Schwarz inequality gives

$$|V_{i,j}|^2 \leq i \sum_{r=0}^{i-1} d_{r,j}^2 \leq N \sum_{r=0}^{N-1} d_{r,j}^2.$$

Using the zero Dirichlet values in (3.3), and $hN = 1$,

$$\|V\|_{2,h}^2 = h^2 \sum_{j=1}^{N-1} \left(\sum_{i=1}^{N-1} V_{i,j}^2 + \tfrac{1}{2} V_{N,j}^2 \right) \le h^2 N^2 S_x(V) = S_x(V).$$

The same pointwise bound with $i = N$ yields $\|V_B\|_{\Gamma,h}^2 = h \sum_j V_{N,j}^2 \le hN S_x(V) = S_x(V)$. Finally (5.3) bounds $S_x(V)$ by the energy. □

Proposition 5.3 (Uniform energy stability)**.** *Under the hypotheses of Proposition 4.6,*

$$\|V\|_{E,h,c} \le \|F\|_{I,h} + \gamma_c \|S\|_{\Gamma,h}, \qquad \gamma_{\mathrm{FO}} = 1, \quad \gamma_{\mathrm{G}} = h/2. \tag{5.8}$$

Proof. Let w_c be the boundary weight in (5.1), so that $\gamma_c = h w_c$. The equation $A_h^c \boldsymbol{V} = \boldsymbol{d}$, where $\boldsymbol{d}$ has entries F and S, gives

$$\|V\|_{E,h,c}^2 = h^2 \sum_{\mathcal{I}_h} V_{i,j} F_{i,j} + h^2 w_c \sum_{j=1}^{N-1} V_{N,j} S_j.$$

Cauchy–Schwarz and Lemma 5.2 imply

$$\begin{aligned} \|V\|_{E,h,c}^2 &\le \|V_I\|_{I,h} \|F\|_{I,h} + h w_c \|V_B\|_{\Gamma,h} \|S\|_{\Gamma,h} \\ &\le \|V\|_{E,h,c} \big(\|F\|_{I,h} + \gamma_c \|S\|_{\Gamma,h} \big). \end{aligned}$$

Here $\|V_I\|_{I,h} \le \|V\|_{2,h} \le \|V\|_{E,h,c}$. If the energy norm is zero the claim is immediate; otherwise division by it proves (5.8). □

Corollary 5.4 (Energy and discrete L^2 convergence)**.** *Under Assumption 2.1, the discretisation errors satisfy*

$$\|e_h^{\mathrm{FO}}\|_{2,h} \le \|e_h^{\mathrm{FO}}\|_{E,h,\mathrm{FO}} \le \tfrac{M_4}{6} h^2 + \tfrac{M_2}{2} h, \tag{5.9}$$

$$\|e_h^{\mathrm{G}}\|_{2,h} \le \|e_h^{\mathrm{G}}\|_{E,h,\mathrm{G}} \le \tfrac{M_4 + M_3}{6} h^2 + \tfrac{M_4}{24} h^3. \tag{5.10}$$

Thus both norms have upper orders one (FO) and two (G).

Proof. Apply (5.8) to the error equation with $F = -\tau_I$ and $S = -\tau_B^c$. Since $h^2(N-1)^2 \le 1$ and $h(N-1) \le 1$, the restricted norms satisfy $\|\tau_I\|_{I,h} \le \|\tau_I\|_{\infty,I}$ and $\|\tau_B^c\|_{\Gamma,h} \le \|\tau_B^c\|_{\infty,B}$. Insert (4.6). The boundary factors are 1 for FO and $h/2$ for G, giving the two energy bounds exactly as stated. Lemma 5.2 supplies the L^2 bounds. □

The energy norm depends on the selected boundary closure. It provides a further theoretical and numerical check, while the common nodal maximum norm remains the primary target for comparing computational cost.

5.3 Solver implications of the proved matrix structure

The raw matrices are nonsymmetric and invertible. We use a sparse direct baseline through SciPy's `spsolve` interface [6], with CSC input and `use_umfpack=False`, explicitly fixing the SuperLU backend. Each call includes a fresh factorisation; no factor is reused between timed repetitions. Multiplication by the invertible positive W_h^c gives the equivalent system $K_h^c \boldsymbol{U}_h^c = W_h^c \boldsymbol{b}_h^c$. Proposition 5.1 proves the symmetry and positive definiteness required for CG [7]. Its diagonal is positive because $(K_h^c)_{kk} = \boldsymbol{e}_k^T K_h^c \boldsymbol{e}_k > 0$. Consequently, the Jacobi inverse $\mathrm{diag}(K_h^c)^{-1}$ is an admissible positive definite preconditioner. The baseline and cross-checks below test these solver choices; the structural proof alone makes no runtime claim.

6 Algorithm and verification protocol

6.1 Assembly, library interfaces and reproducible outputs

The project-specific code implements grid ownership, the ordering (3.2), both boundary closures, exact Dirichlet elimination, error norms, componentwise residuals, row weights and the edge-sum energy. Sparse entries are collected in COO and converted to CSR; the direct solve converts the matrix to CSC. The present square-grid assembly uses explicit stencil loops for clarity. Array operations evaluate norms and most diagnostics; the measured assembly cost includes the loops. NumPy supplies array operations and a dense small-system solve; SciPy supplies sparse storage, SuperLU, GMRES, ILU, CG and the symmetric eigenvalue routine. Matplotlib supplies figures. These interfaces are delegated numerical services, rather than separately implemented solvers.

1. **Input:** $N \geq 2$, closure $c \in \{\mathrm{FO}, \mathrm{G}\}$, $\beta \in \{0, 2\}$, and independent functions u^*, f, s_β. Set $h = 1/N$ and enumerate $\mathcal{U}_h$ by (3.2).
2. For each interior node, assemble (3.5). For each open right-edge node, assemble (3.6) or (3.10). Do not assign a derivative row to a corner.
3. For each known Dirichlet neighbour with coefficient a_{kD}, add $-a_{kD}g_D$ to the right-hand side. Store only unknown columns.
4. Convert COO entries to CSR, sum duplicate entries, and check the dimensions and finiteness of $A_h^c, \boldsymbol{b}_h^c$.
5. Solve $A_h^c \widetilde{\boldsymbol{U}}_h^c = \boldsymbol{b}_h^c$ using the fixed sparse direct backend. Reconstruct all grid values, retaining exact Dirichlet values.
6. Evaluate the global norms (3.3), restricted boundary/interior errors, residuals (6.1), matrix asymmetry and, where required, the weighted energy error.
7. **Output:** solution, diagnostics and measured times. The manufactured error is evaluated against u^*, not inferred from the residual.

Algorithm 6.1: A mixed-boundary finite difference solve and its diagnostics

The script `code/run_all.py` executes the four configurations, the four-mesh study, fixed-target cost study and solver cross-checks. It generates CSV data in `data/`, LaTeX tables in `generated/`, and figures in `figures/`. The report author can include the generated tables directly, avoiding manual re-entry of numerical values.

6.2 Checks that test the equations and checks that test the solve

The executable test suite has ten test methods. It checks dimensions, signs, finite entries, invalid options, corner ownership, exact rows, manufactured defects, observed rates and the weighted energy identity. We distinguish the following independent checks.

Hand-derived boundary row. At $N = 4$, $h = 1/4$, $j = 1$, the ghost row has diagonal $64 + 8\beta$, west coefficient -32 and retained north coefficient -16. After eliminating the bottom-right corner, its right-hand side is $(6 + 8\beta)e^{1.25} + 16e$. These values check the inward factor, outward-normal sign and corner contribution separately.

Manufactured functions with different polynomial structure. For $u = 1 + x + y$, both closures have zero defect in exact arithmetic. For $u = 1 + x^2 + y^2$, the ghost defect is zero and the FO boundary defect is $-h$. An additional check uses

$$u = 1 + x^3 + y^4, \qquad f = -6x - 12y^2, \qquad s_\beta = 3 + \beta(2 + y^4).$$

Direct polynomial substitution predicts

$$\tau_I = -2h^2, \qquad \tau_B^{\mathrm{FO}} = -3h + h^2, \qquad \tau_B^{\mathrm{G}} = 2h - 2h^2.$$

These identities are checked for both boundary variants at $N = 2, 3, 4, 6$, using data derived without applying the discrete matrix. They exercise nonzero forcing as well as boundary derivatives.

A separately assembled uneliminated system. On the same four small grids, an additional dense system retains every Dirichlet variable and, for G, every ghost variable. It imposes explicit Dirichlet identity equations, the original five-point PDE and centred boundary equations. Its physical-node solution is compared with the reduced sparse solve in all sixteen configurations. Agreement is required within absolute and relative tolerances 10^{-11}. This checks elimination and assembly through a different system layout.

Same-system solver checks. A dense solve of the same reduced matrix and the GMRES/CG comparisons test linear algebra. Agreement alone would not establish that the matrix discretises the intended PDE. The hand rows, independent data and uneliminated construction supply the complementary equation checks. Automated algebraic checks do not constitute a human audit sign-off.

6.3 Residuals, tolerance sweep and acceptance criteria

For a computed vector $\widetilde{\boldsymbol{U}}$, set $\boldsymbol{r} = \boldsymbol{b} - A\widetilde{\boldsymbol{U}}$ and report

$$\eta_{\mathrm{cw}} = \max_k \frac{|r_k|}{(|A||\widetilde{\boldsymbol{U}}| + |\boldsymbol{b}|)_k}, \qquad \eta_{A,\infty} = \frac{\|\boldsymbol{r}\|_\infty}{\|A\|_\infty \|\widetilde{\boldsymbol{U}}\|_\infty + \|\boldsymbol{b}\|_\infty}. \tag{6.1}$$

A zero denominator with zero numerator contributes zero. The componentwise measure is unchanged by positive diagonal row scaling, so it compares the different raw row scales meaningfully. Neither residual is a discretisation error or a forward-error bound. The baseline is direct and has no user-set stopping tolerance.

At $N = 64$, we independently cross-check its solution using GMRES [8]. Let $D_{kk} = (\sum_j |a_{kj}|)^{-1}$; GMRES solves $DAU = Db$ with ILU [9], drop tolerance 10^{-4}, fill factor 20, restart 50, at most 1000 restart cycles, zero initial vector, zero absolute tolerance and relative tolerances $10^{-8}, 10^{-10}, 10^{-12}$. The residual stopping test uses the scaled system's Euclidean norm; the original-system diagnostics in (6.1) are evaluated separately. D is not the analytical symmetriser W_h^c.

The acceptance measure is

$$\rho_{\mathrm{alg}} = \frac{\|U_{\mathrm{iterative}} - U_{\mathrm{direct}}\|_\infty}{\|U_{\mathrm{direct}} - u_h^*\|_\infty} < 10^{-2}. \tag{6.2}$$

The boundary-mean quantity $Q_h = (N-1)^{-1} \sum_{j=1}^{N-1} U_{N,j}$ is monitored as well. Agreement at tighter tolerances supports negligible solve error on this representative grid; it is not a general proof for arbitrary systems. Weighted CG uses $K_h^c U = W_h^c b$, Jacobi preconditioning, zero initial vector, rtol 10^{-10}, atol zero and at most 10000 iterations. All iterative exit codes and residuals are recorded.

For the energy identity, a seeded random vector is tested against the independent edge sums, with relative discrepancy below 10^{-12}. At $N = 32$, `eigsh` with `which='SA'` and tolerance 10^{-10} estimates the smallest eigenvalue [10]. This is a finite numerical corroboration of Proposition 5.1, not its proof.

7 Numerical experiments and interpretation

7.1 Four-mesh protocol and global errors

Both boundary variants and both closures use the same data, ordering and Dirichlet ownership. The required refinement family is $N = 16, 32, 64, 128$, or $h = 1/16, 1/32, 1/64, 1/128$. For a computed full-grid error $e = \widetilde{U} - u_h^*$, we report $E_{2,h} = \|e\|_{2,h}$ and $E_{\infty,h} = \|e\|_{\infty,h}$ using exactly (3.3). At consecutive levels,

$$p_m(N) = \frac{\log(E_{m,2h}/E_{m,h})}{\log 2}, \qquad m = 2, \infty.$$

After inspecting these pairwise trends, the final three consecutive levels define the fitted range. We fit $\log E = p \log h + C$ by least squares and report p and R^2. A good straight-line fit over three points is evidence for an observed rate on that range, not a proof of an asymptotic expansion.

The bundled data were generated on Windows 11 with Python 3.13.15, NumPy 2.3.5, SciPy 1.17.0 and Matplotlib 3.10.8. The version and timing settings are recorded in `data/run_metadata.json`. Reproduction on another machine may change timings and the last few residual digits.

N	h	$E_{2,h}$	p_2	$E_{\infty,h}$	p_∞	η_{cw}
16	1/16	1.786×10^{-2}	–	5.832×10^{-2}	–	1.95×10^{-14}
32	1/32	8.515×10^{-3}	1.069	2.815×10^{-2}	1.051	2.75×10^{-14}
64	1/64	4.161×10^{-3}	1.033	1.383×10^{-2}	1.025	7.32×10^{-14}
128	1/128	2.058×10^{-3}	1.016	6.853×10^{-3}	1.013	2.17×10^{-13}

Table 1: Neumann data, FO closure: four-mesh errors and pairwise orders. The direct SuperLU solver has no stopping tolerance; η_{cw} is the original-system componentwise backward residual.

N	h	$E_{2,h}$	p_2	$E_{\infty,h}$	p_∞	η_{cw}
16	1/16	2.206×10^{-4}	–	8.611×10^{-4}	–	2.41×10^{-16}
32	1/32	5.499×10^{-5}	2.004	2.163×10^{-4}	1.993	3.78×10^{-16}
64	1/64	1.374×10^{-5}	2.001	5.412×10^{-5}	1.998	3.37×10^{-16}
128	1/128	3.434×10^{-6}	2.000	1.354×10^{-5}	1.999	4.44×10^{-16}

Table 2: Neumann data, G closure: four-mesh errors and pairwise orders. The direct SuperLU solver has no stopping tolerance; η_{cw} is the original-system componentwise backward residual.

N	h	$E_{2,h}$	p_2	$E_{\infty,h}$	p_∞	η_{cw}
16	1/16	1.059×10^{-2}	–	3.430×10^{-2}	–	4.89×10^{-15}
32	1/32	5.143×10^{-3}	1.042	1.685×10^{-2}	1.026	3.50×10^{-14}
64	1/64	2.536×10^{-3}	1.020	8.355×10^{-3}	1.012	9.16×10^{-14}
128	1/128	1.259×10^{-3}	1.010	4.159×10^{-3}	1.007	2.39×10^{-13}

Table 3: Robin data, FO closure: four-mesh errors and pairwise orders. The direct SuperLU solver has no stopping tolerance; η_{cw} is the original-system componentwise backward residual.

N	h	$E_{2,h}$	p_2	$E_{\infty,h}$	p_∞	η_{cw}
16	1/16	1.273×10^{-4}	–	5.307×10^{-4}	–	2.27×10^{-16}
32	1/32	3.162×10^{-5}	2.009	1.332×10^{-4}	1.994	2.81×10^{-16}
64	1/64	7.892×10^{-6}	2.003	3.335×10^{-5}	1.998	4.60×10^{-16}
128	1/128	1.972×10^{-6}	2.001	8.340×10^{-6}	2.000	4.16×10^{-16}

Table 4: Robin data, G closure: four-mesh errors and pairwise orders. The direct SuperLU solver has no stopping tolerance; η_{cw} is the original-system componentwise backward residual.

Tables 1–4 show decreasing errors on every refinement. FO pairwise orders approach one from above in both global norms, while G orders are close to two. The maximum componentwise direct-solve residual across these sixteen runs is 2.39×10^{-13}. Its small size is supplemented by the independent solve checks below; it is not used alone to establish correct boundary discretisation.

Boundary	Closure	p_2	R_2^2	p_∞	R_∞^2	p_B	p_I
Neumann	FO	1.025	0.999977	1.019	0.999988	1.019	0.971
Neumann	G	2.001	1.000000	1.999	1.000000	1.999	1.933
Robin	FO	1.015	0.999991	1.009	0.999997	1.009	0.960
Robin	G	2.002	1.000000	1.999	1.000000	1.999	1.918

Table 5: Least-squares fits on the identified fine-mesh range $N = 32, 64, 128$. Pairwise orders are retained in the preceding refinement tables; R^2 measures fit quality rather than a confidence interval.

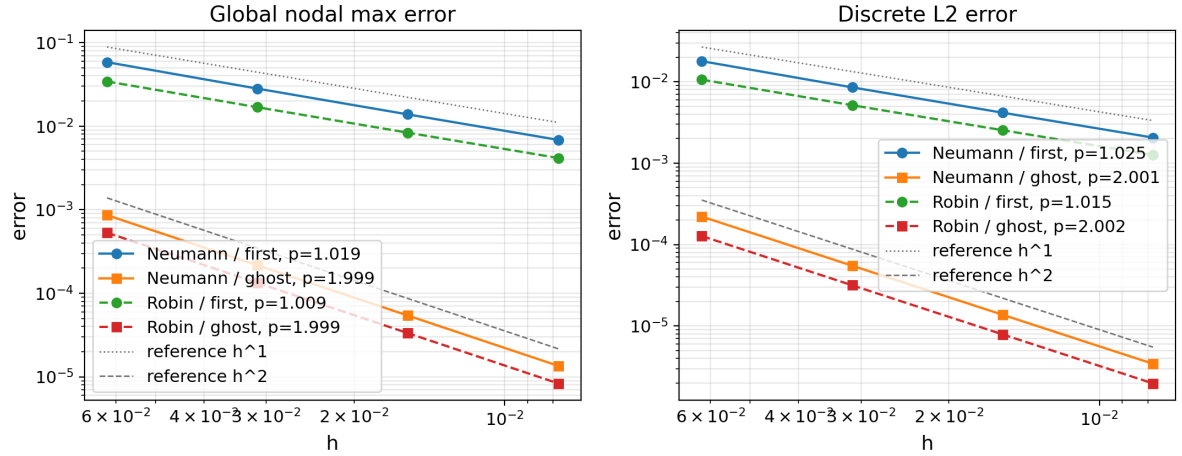


Figure 2: Global maximum and trapezoidal discrete L^2 refinement errors. Legends give the fitted slopes on $N = 32, 64, 128$; dotted/dashed reference curves have slopes one and two and arbitrary vertical placement. Errors approach the orders predicted by Theorem 4.7 and Corollary 5.4.

Table 5 distinguishes a theorem's upper order from a measured slope. The FO maximum-error rate is also supported by the independent lower bound in Theorem 4.8. For G, observed second order agrees with the proved upper bound; no mathematical lower bound for its error is inferred from these four meshes.

7.2 Boundary error, interior error and finite-grid deviations

We evaluate $E_{\infty,B} = \max_{1 \le j < N} |e_{N,j}|$ and $E_{\infty,I} = \max_{(i,j) \in \mathcal{I}_h} |e_{i,j}|$ separately. The CSV data also retain the restricted L^2 errors defined by (3.4). This classifies errors by where they are

measured. It is not an additive separation of the effects of boundary and interior truncation defects, which are coupled by the elliptic solve.

Closure	N	$E_{\infty,B}$	p_B	$E_{\infty,I}$	p_I
FO	16	5.832×10^{-2}	–	4.863×10^{-2}	–
FO	32	2.815×10^{-2}	1.051	2.574×10^{-2}	0.918
FO	64	1.383×10^{-2}	1.025	1.322×10^{-2}	0.961
FO	128	6.853×10^{-3}	1.013	6.701×10^{-3}	0.981
G	16	8.611×10^{-4}	–	6.706×10^{-4}	–
G	32	2.163×10^{-4}	1.993	1.914×10^{-4}	1.809
G	64	5.412×10^{-5}	1.998	5.096×10^{-5}	1.909
G	128	1.354×10^{-5}	1.999	1.314×10^{-5}	1.956

Table 6: Neumann data: error maxima restricted to the open right edge and the interior. This is a partition by node location, not an additive decomposition by error source.

Closure	N	$E_{\infty,B}$	p_B	$E_{\infty,I}$	p_I
FO	16	3.430×10^{-2}	–	2.854×10^{-2}	–
FO	32	1.685×10^{-2}	1.026	1.537×10^{-2}	0.893
FO	64	8.355×10^{-3}	1.012	7.978×10^{-3}	0.946
FO	128	4.159×10^{-3}	1.007	4.064×10^{-3}	0.973
G	16	5.307×10^{-4}	–	3.928×10^{-4}	–
G	32	1.332×10^{-4}	1.994	1.149×10^{-4}	1.774
G	64	3.335×10^{-5}	1.998	3.098×10^{-5}	1.890
G	128	8.340×10^{-6}	2.000	8.039×10^{-6}	1.946

Table 7: Robin data: error maxima restricted to the open right edge and the interior. This is a partition by node location, not an additive decomposition by error source.

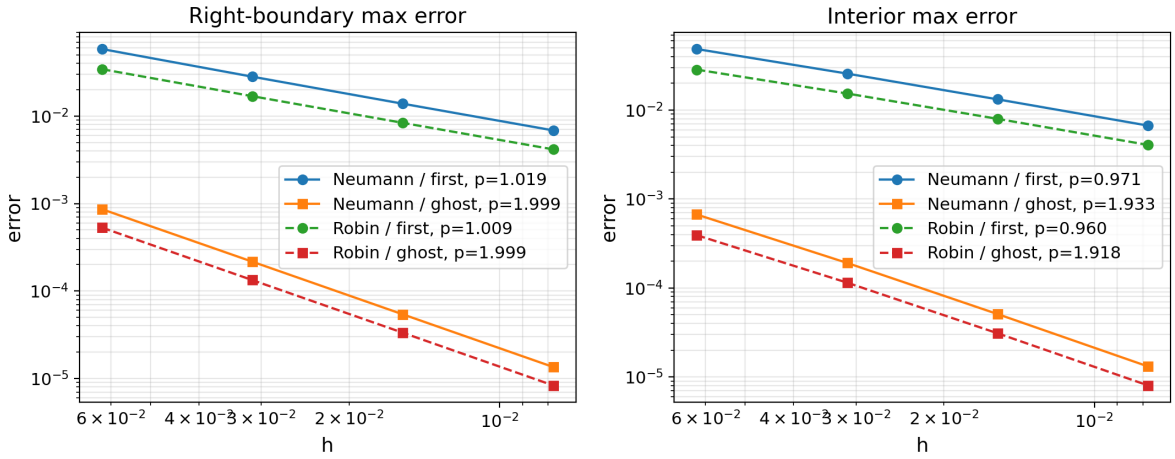


Figure 3: Error maxima on the open right edge and in the interior. The displayed slopes use the same finest-three-level fits. FO boundary error limits global accuracy, and the associated error extends into the interior. The spatial classification does not attribute interior error exclusively to the interior stencil.

For all tested configurations the global maximum equals the right-edge maximum, and the interior maximum is smaller but of comparable size. The interior FO fits approach one, consistent

with a boundary defect affecting the full field. Thus a second-order interior stencil does not give a second-order global solution under the FO closure.

The G interior maximum fits are [define value here] and 1.918, below the approximately two global fits. Their last pairwise orders increase to approximately 1.956 and 1.946. An additional location diagnostic, `data/error_locations.csv`, finds the fine-grid interior maxima on column $i = N - 1$. This sampled column moves towards $x = 1$ as h decreases. The interior/global error ratios also move towards one. These observations indicate finite-grid and maximising-location effects; they do not identify a separate asymptotic leading term. The measured deviations are compatible with the uniform $O(h^2)$ upper bound and are not evidence of tolerance contamination, given the much smaller solver differences below.

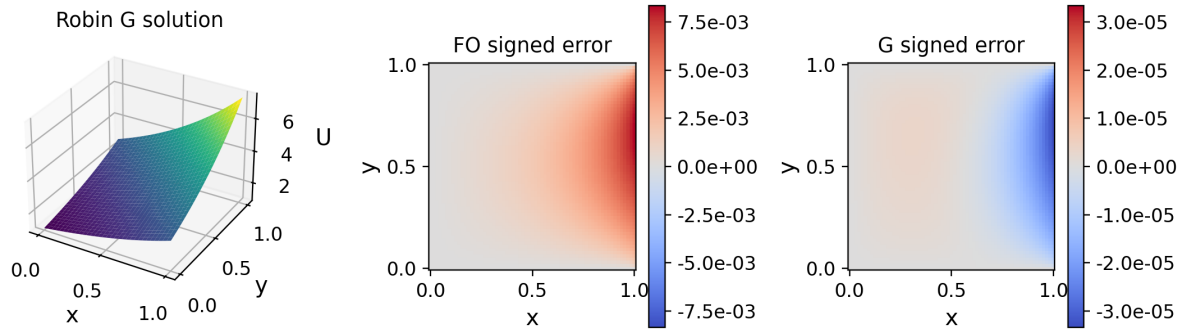


Figure 4: Representative Robin solution and signed errors at $N = 64$. The surface is the G solution. The FO and G error maps use separate, explicit colour scales centred at zero; their colours alone should not be compared as equal error magnitudes. All three Dirichlet edges and four corners have zero imposed-data error.

7.3 Raw matrices, weighted structure and sparsity

At $N = 32$, $n_{\text{dof}} = 992$, and the exact counts in Proposition 3.2 give 4774 FO and 4834 G nonzeros. The spy plots show the extra right-edge tangential couplings in G. They do not themselves establish symmetry; the Frobenius diagnostic quantifies it.

Boundary	Closure	δ_A	δ_K	$\lambda_{\min}(K)$	Energy defect
Neumann	FO	0.05515	0	11.6098	0
Neumann	G	0.05595	0	11.9409	1.19×10^{-16}
Robin	FO	0.05515	0	14.7168	1.14×10^{-16}
Robin	G	0.05590	0	14.8745	2.33×10^{-16}

Table 8: Raw and weighted matrix checks at $N = 32$: $\delta_A = \|A - A^T\|_F / \|A\|_F$ and $\delta_K = \|K - K^T\|_F / \|K\|_F$. Eigenvalues are numerical estimates; positive definiteness for all stated grids is proved in Proposition 5.1.

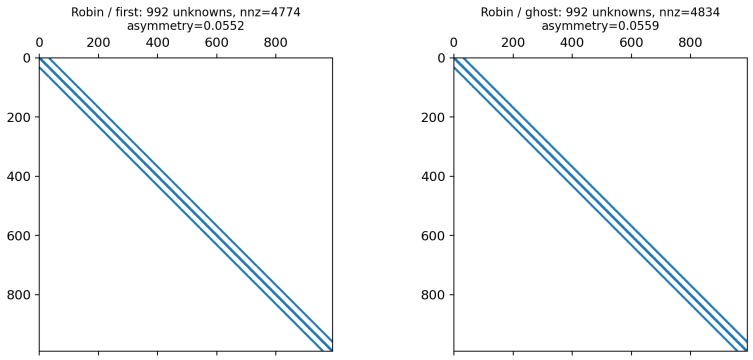


Figure 5: Representative Robin raw matrices at $N = 32$. The sparsity patterns retain nearest-neighbour couplings, while ghost elimination changes the inward coefficient and adds tangential boundary couplings. Both raw matrices are nonsymmetric, as confirmed by the labelled Frobenius ratios.

Table 8 gives nonzero raw asymmetry for all four configurations, exactly zero measured weighted asymmetry on this binary grid, and positive smallest-eigenvalue estimates. The edge-sum identity agrees with the quadratic form at roundoff. These checks corroborate the particular implementation of W_h^c, while Proposition 5.1 supplies the general proof.

7.4 Weighted-energy convergence

For the error restricted to unknowns we compute

$$E_{E,h,c} = \left(h^2 \boldsymbol{e}^T K_h^c \boldsymbol{e}\right)^{1/2}.$$

Each quadratic form is independently evaluated through the edge sums. The largest relative discrepancy over the refinement runs is [define value here]. The four-mesh data remain in `data/energy_convergence.csv`.

Boundary	Closure	$E_{E,h}$ at $N = 128$	Fitted p_E	R_E^2
Neumann	FO	9.581×10^{-3}	1.010	0.999997
Neumann	G	2.018×10^{-5}	1.997	1.000000
Robin	FO	7.596×10^{-3}	1.005	1.000000
Robin	G	1.647×10^{-5}	1.996	1.000000

Table 9: Measured energy errors and fine-mesh fits on $N = 32, 64, 128$. The energy norm is closure dependent; it is not the shared cost-comparison norm.

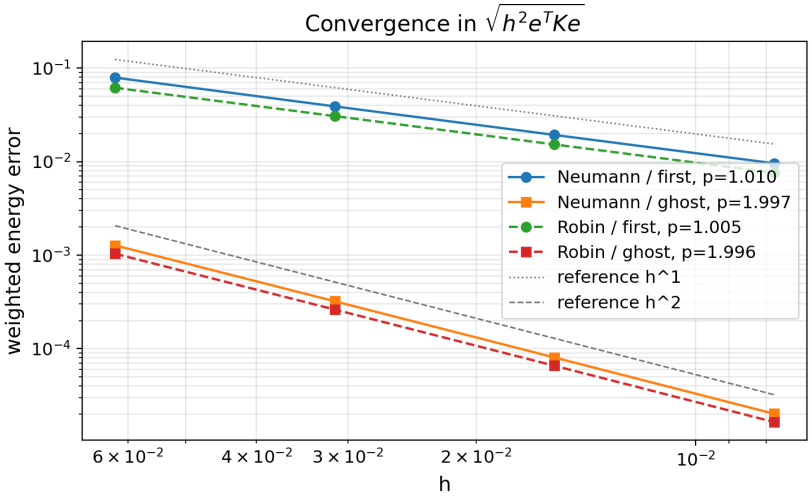


Figure 6: Convergence in the proved closure-dependent energy norm. Reference slopes one and two and fitted slopes on $N = 32, 64, 128$ are shown. FO and G approach the respective upper orders in Corollary 5.4.

The energy fits approach one and two, respectively. This additional norm checks the theory beyond nodal errors, but the two closures use different energy norms. The common nodal maximum error, rather than these closure-dependent energies, therefore defines the cost target.

7.5 Algebraic-error separation and weighted CG

Table 10 implements the tolerance sweep specified in Section 6. Every run returns status zero. Even at rtol 10^{-8} the largest ratio in (6.2) is 3.67×10^{-5}, below 10^{-2}; at 10^{-12} it is 4.30×10^{-8}. At the tightest tolerance, the largest change in the boundary-mean quantity relative to the direct solution is 1.75×10^{-12}. The residual, solution difference and quantity of interest thus provide separate evidence for the linear solve.

Boundary	Closure	rtol	Iterations	η_{cw}	ρ_{alg}
Neumann	FO	1×10^{-8}	4	5.92×10^{-11}	6.69×10^{-8}
Neumann	FO	1×10^{-10}	5	1.37×10^{-13}	2.55×10^{-10}
Neumann	FO	1×10^{-12}	5	9.59×10^{-14}	2.91×10^{-10}
Neumann	G	1×10^{-8}	4	7.09×10^{-11}	2.20×10^{-5}
Neumann	G	1×10^{-10}	5	1.95×10^{-13}	8.92×10^{-8}
Neumann	G	1×10^{-12}	5	6.52×10^{-14}	3.55×10^{-8}
Robin	FO	1×10^{-8}	4	6.48×10^{-11}	1.17×10^{-7}
Robin	FO	1×10^{-10}	5	2.14×10^{-13}	5.19×10^{-10}
Robin	FO	1×10^{-12}	5	1.20×10^{-13}	3.66×10^{-10}
Robin	G	1×10^{-8}	4	9.19×10^{-11}	3.67×10^{-5}
Robin	G	1×10^{-10}	5	1.21×10^{-13}	1.07×10^{-7}
Robin	G	1×10^{-12}	5	9.91×10^{-14}	4.30×10^{-8}

Table 10: GMRES–ILU cross-check at $N = 64$, with atol $= 0$. Residuals are evaluated on the original raw system. The ratio ρ_{alg} is the iterative/direct maximum difference divided by the measured direct-solution error; the acceptance target is 10^{-2}.

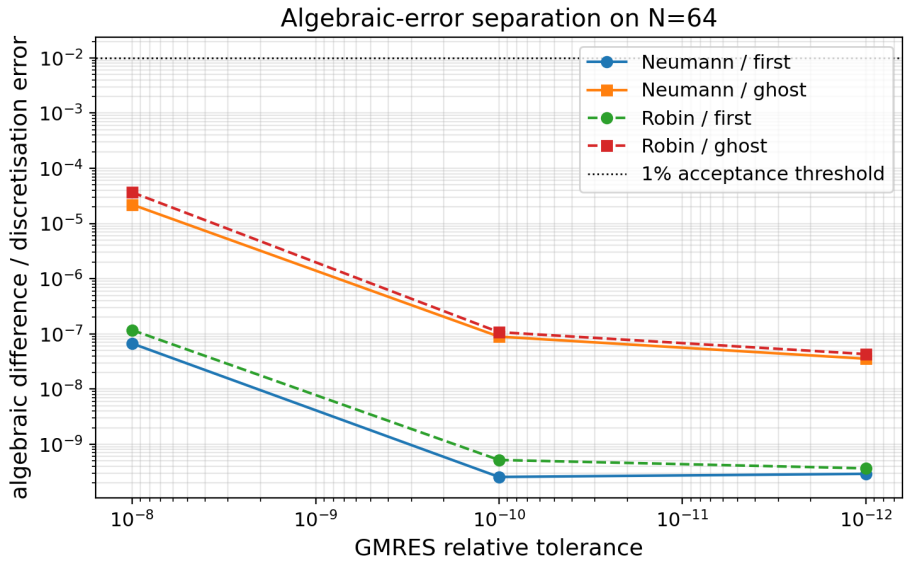


Figure 7: GMRES/direct differences relative to discretisation errors at $N = 64$. All ratios lie below the one-percent acceptance line. The flattening at tight tolerances is consistent with a solve-difference floor and is not interpreted as a new discretisation rate.

For the weighted CG check, the reported normwise weighted residual is

$$\eta_{K,\infty} = \frac{\|Wb - K\widetilde{U}\|_\infty}{\|K\|_\infty \|\widetilde{U}\|_\infty + \|Wb\|_\infty}.$$

Table 11 also gives the original-system η_{cw}, so differently scaled residuals are not conflated. All four CG solves return status zero and require 280–289 iterations. The largest CG/direct difference ratio is 3.06×10^{-5}. This verifies the implementation of the proved eligible system. It does not establish that CG is faster than direct factorisation.

Boundary	Closure	Iterations	$\eta_{K,\infty}$	η_{cw}	ρ_{alg}
Neumann	FO	289	4.30×10^{-12}	1.98×10^{-11}	6.82×10^{-8}
Neumann	G	287	3.79×10^{-12}	1.66×10^{-11}	1.63×10^{-5}
Robin	FO	280	5.27×10^{-12}	2.57×10^{-11}	8.97×10^{-8}
Robin	G	280	3.84×10^{-12}	1.87×10^{-11}	3.06×10^{-5}

Table 11: Jacobi-preconditioned CG on $KU = Wb$ at $N = 64$, rtol 10^{-10}, atol zero. All return status zero. The two residual columns distinguish weighted normwise and original componentwise measures.

These comparisons use independent linear solvers for the same system; they do not replace the assembly checks in Section 6. The tolerance sweep is conducted on the required representative mesh rather than claimed as a forward-error proof for every possible refinement.

7.6 Measured cost at fixed accuracy thresholds

The cost measure is wall-clock time for assembly followed by a fresh sparse direct solve, including COO/CSR/CSC conversion, ordering, factorisation and triangular solves. Imports, a warm-up run, post-solve error/residual diagnostics, independent tests, Krylov and eigenvalue cross-checks, plotting and file I/O are excluded. After one warm-up, five complete totals T_k are measured with `perf_counter`; the reported cost is $\mathrm{median}_k\, T_k$. Separate assembly and solve medians are

retained in CSV, but their sum is not substituted for this total median. The displayed minimum–maximum is sample variability, not a confidence interval. No factorisation or preconditioner is reused.

The fixed targets are $E_{\infty,h} \leq 10^{-2}$ and $E_{\infty,h} \leq 5 \times 10^{-3}$, for each boundary variant separately. Candidate grids are $N \in \{4, 8, 16, 32, 64, 128\}$; the two extra coarse levels serve this cost study only and do not enter convergence fits. The prescribed selection rule chooses the coarsest tested grid meeting the target. If none does, the target is recorded as unattained; no unmeasured runtime is extrapolated. This is a comparison at a common error bound, not at identical achieved error. In particular, G may oversatisfy a threshold, and the selected grid is not asserted to be a global runtime optimum over all possible N.

Boundary	Target	Closure	N	$E_{\infty,h}$	Median (ms)	Range (ms)
Neumann	1×10^{-2}	FO	128	6.853×10^{-3}	100.000	98.886–106.615
Neumann	1×10^{-2}	G	8	3.377×10^{-3}	1.119	0.888–1.629
Neumann	5×10^{-3}	FO	–	not reached	–	–
Neumann	5×10^{-3}	G	8	3.377×10^{-3}	1.119	0.888–1.629
Robin	1×10^{-2}	FO	64	8.355×10^{-3}	22.096	20.647–22.663
Robin	1×10^{-2}	G	4	7.326×10^{-3}	0.373	0.368–0.382
Robin	5×10^{-3}	FO	128	4.159×10^{-3}	100.318	98.051–102.029
Robin	5×10^{-3}	G	8	2.066×10^{-3}	0.822	0.815–1.339

Table 12: Cost at fixed nodal maximum-error thresholds, selecting the coarsest successful grid in $\{4, 8, 16, 32, 64, 128\}$. Each range is the minimum–maximum of five measured assembly-plus-solve totals after one warm-up, not a confidence interval. A missing entry means the target was not attained within the candidate set.

For Neumann data at target 1×10^{-2}, the selected FO/G grids are $N = 128$ and $N = 8$; their measured median cost ratio $T_{\mathrm{FO}}/T_{\mathrm{G}}$ is 89.33.

For Neumann data at target 5×10^{-3}, FO does not meet the target on any tested grid up to $N = 128$, while G meets it at $N = 8$. No finite FO/G cost ratio is reported for this unattained case.

For Robin data at target 1×10^{-2}, the selected FO/G grids are $N = 64$ and $N = 4$; their measured median cost ratio $T_{\mathrm{FO}}/T_{\mathrm{G}}$ is 59.19.

For Robin data at target 5×10^{-3}, the selected FO/G grids are $N = 128$ and $N = 8$; their measured median cost ratio $T_{\mathrm{FO}}/T_{\mathrm{G}}$ is 122.10.

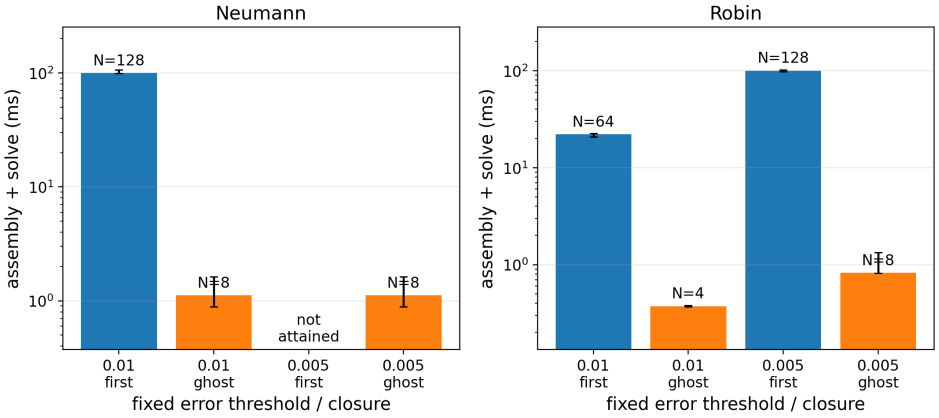


Figure 8: Measured cost under the same error thresholds within each boundary variant. The vertical axis is logarithmic. Error bars show the five-sample minimum–maximum; labels identify the selected grids. Neumann FO fails the tighter threshold within the tested range, so no cost bar is fabricated for that case.

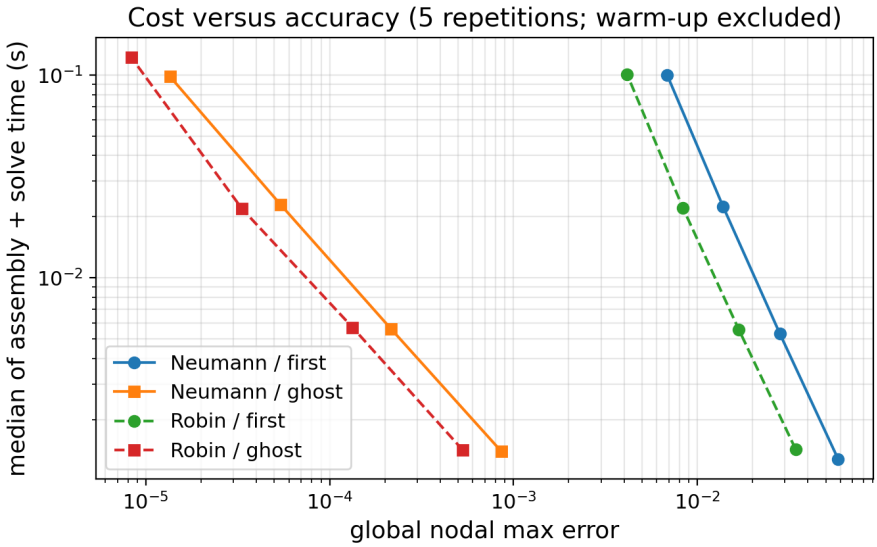


Figure 9: The full error–cost trajectories on the four refinement grids. Cost is the median of measured assembly-plus-solve totals, including factorisation. This provides context; quantitative fixed-target comparisons use Table 12, not a visual comparison of unmatched points.

For the successful comparisons on this implementation, G reaches the thresholds on much smaller systems and has lower measured median cost. That is a measured conclusion for these data, candidate grids and machine, rather than a hardware-independent complexity claim. The unattained Neumann FO case is part of the result. Timings are regenerated by the driver; their ratios should not be transplanted to another machine. No direct/CG runtime crossover is claimed from the single-grid CG cross-check.

8 Conclusions and limitations

The right-edge equation controls the accuracy retained by the mixed Poisson discretisation. The two raw boundary defects are $O(h)$, but the comparison estimate weights the G defect by h. Together with the $O(h^2)$ interior defect, this yields the FO $O(h)$ and G $O(h^2)$ maximum-norm bounds. For e^{x+y}, the separate FO lower bound proves that first order is sharp in the boundary and global maximum norms. The weighted energy identity supplies a second stability route and discrete L^2/energy bounds for the actual boundary and corner layout.

Four-mesh experiments for both Neumann and Robin data reproduce the predicted approximately first- and second-order behaviour. The interior maximum-error fits exhibit visible finite-grid deviations that are reported rather than rounded to the theoretical orders. Independent boundary rows, different manufactured functions and an uneliminated system check test the equations. Direct/GMRES comparisons and weighted CG show solve differences much smaller than the measured discretisation error on the representative grid. The implemented positive row weights recover the proved symmetric positive definite structure without asserting symmetry of the raw systems.

At fixed nodal-error thresholds, the measured G costs are lower in the successful comparisons because substantially coarser tested grids meet the same bounds. Neumann FO does not attain the tighter target within $N \leq 128$; no cost ratio is assigned to that case. Cost includes fresh factorisation and is measured as complete assembly-plus-solve totals. The finite candidate set, coarsest-successful-grid policy and timing variability limit the scope of this efficiency conclusion.

The mathematical assumptions remain essential: bounded derivatives, fixed nonnegative β, exact imposed Dirichlet values and the stated square geometry. Smoothness is satisfied by the manufactured function and is not proved for arbitrary mixed data or corner-singular problems. The G theorem is an upper bound without a matching lower bound. Floating-point diagnostics and representative solver cross-checks are evidence for the reported computations, not universal forward-error certificates. The energy norms depend on the closure; common nodal errors define the cost targets. Finally, this is verification on known solutions rather than experimental validation of a particular device.

The analysis and computations jointly support the conclusion that the ghost-eliminated formulation preserves second-order accuracy for this smooth benchmark, whereas the backward derivative formulation has a provably first-order maximum error. The distinction follows from the complete boundary formulation and its stability, not from the interior stencil alone.

References

[1] R. J. LeVeque (2007). *Finite Difference Methods for Ordinary and Partial Differential Equations: Steady-State and Time-Dependent Problems*. Society for Industrial and Applied Mathematics, Philadelphia. Author's University of Washington book page: https://faculty.washington.edu/rjl/fdmbook/index.html.

[2] K. Salari and P. Knupp (2000). *Code Verification by the Method of Manufactured Solutions*. Sandia National Laboratories, SAND2000-1444. DOI: 10.2172/759450. Official Sandia publication record: https://www.sandia.gov/research/publications/details/code-verification-by-the-method-of-manufactured-solutions-2000-06-01/.

[3] J. Schöberl (n.d.). *The weak formulation of the Poisson equation*. An Interactive Introduction to the Finite Element Method, Chapter 17. Author's online textbook: https://jschoeberl.github.io/iFEM/sobolevspaces/preciseweak.html.

[4] P.-O. Persson (n.d.). *Finite Difference Methods for PDEs*. Math 228B, University of California, Berkeley. Official university course notes: https://persson.berkeley.edu/math228b/html/finite_difference_methods_for_PDEs.html.

[5] A. Berman and R. J. Plemmons (1994). *Nonnegative Matrices in the Mathematical Sciences*. Classics in Applied Mathematics 9, Society for Industrial and Applied Mathematics, Philadelphia. Chapter 6, *M-Matrices*, pp. 132–164. DOI: 10.1137/1.9781611971262. Official SIAM chapter record: https://epubs.siam.org/doi/10.1137/1.9781611971262.ch6.

[6] SciPy developers (n.d.). `scipy.sparse.linalg.spsolve`. Official SciPy API documentation: https://docs.scipy.org/doc/scipy/reference/generated/scipy.sparse.linalg.spsolve.html.

[7] SciPy developers (n.d.). `scipy.sparse.linalg.cg`. Official SciPy API documentation: https://docs.scipy.org/doc/scipy/reference/generated/scipy.sparse.linalg.cg.html.

[8] SciPy developers (n.d.). `scipy.sparse.linalg.gmres`. Official SciPy API documentation: https://docs.scipy.org/doc/scipy/reference/generated/scipy.sparse.linalg.gmres.html.

[9] SciPy developers (n.d.). `scipy.sparse.linalg.spilu`. Official SciPy API documentation: https://docs.scipy.org/doc/scipy/reference/generated/scipy.sparse.linalg.spilu.html.

[10] SciPy developers (n.d.). `scipy.sparse.linalg.eigsh`. Official SciPy API documentation: https://docs.scipy.org/doc/scipy/reference/generated/scipy.sparse.linalg.eigsh.html.

[11] Publications Technical Group, American Mathematical Society (2017). *Using the amsthm Package*. Version 2.20.3, September 2017. AMS-LaTeX documentation for theorem and proof typesetting. CTAN distribution document (University of Illinois mirror): https://ctan.math.illinois.edu/macros/latex/required/amscls/doc/amsthdoc.pdf.

All public online references accessed 15 October 2026.

A Prospective coordination plan

The role assignment for this draft combines coordination, mathematical theory and report preparation (roles 1, 2 and 4). Implementation and testing/validation remain with their designated owner (roles 3 and 5). The supplied role cards explicitly allow roles to be merged or redistributed. The plan below uses the teaching weeks in the Project 2 Brief; it is a proposed schedule, not a record of completed meetings, repository activity or submissions.

Week	Coordination and theory/report deliverable	Evidence to retain
5	Record Direction 2, the role assignment and the continuous/discrete problem definition. Establish the intended repository layout `code/`, `figures/`, `report/`.	Problem definition, named owners and genuine repository records once created.
6	Complete the introduction, model, boundary-row derivations and uniform consistency analysis. Agree notation with the implementation owner.	Versioned mathematical text and the accepted sign/index conventions.
7	Prepare the mathematical clinic explanation: outward normal, corner ownership, local/global order and raw/weighted symmetry. Coordinate the readiness review with the validation owner.	Actual clinic readiness and weekly records; no readiness result is assumed here.
8	Polish the analytical report and mathematical diagrams. Coordinate the chosen presentation wave; numerical illustrations are integrated only after their data are supplied.	Revised text, labelled figures and the actual seminar materials.
9	Proofread the integrated report and references. Coordinate the final report, code and presentation submission against the course deadline.	Actual release files, reproduction record and submission receipt.

Table 13: Proposed responsibilities for the roles covered by this draft. The Week 8 presentation wave is voluntary; otherwise the brief specifies Week 9. The common deadline is Monday of Week 10 at 23:59. No calendar date is inferred without the teaching-week calendar.

The coordinator should arrange the role-card 15-minute meeting before each Tutorial or during its team-work block, keep the task/owner/due table current, and ensure that every participant has a named task. The intended Week 5 setup includes the repository and Write access for every teammate and the teacher. These are planned coordination actions, rather than claims that access was granted or commits were made. Weekly challenge consensus, milestone and forum records should describe what actually occurred. The presentation allocation is ten minutes plus three minutes of Q&A, as specified by the brief. The mathematical component should be explainable without reading the report, especially the derivation of (3.10) and the stability weight in (4.10).

This manuscript addresses the derivation and problem-specific stability requirements in rubric Dimension A, and supplies analytical depth through the lower bound and energy estimates for Dimension B. It does not establish general research novelty. Clear exposition, equation provenance and limitations support the writing requirements in Dimension E. Implementation evidence, numerical experiments and actual individual records remain necessary for the whole project's assessment. In particular, a written plan is not evidence that a repository was created, a solver ran or an individual contribution occurred. An ICS and the weekly records must be written from genuine activity, with the required git-traceable details; none are invented in this draft.

B Mathematical drafting declaration and source provenance

AI assistance was used to organise the supplied requirements, draft the English exposition and LaTeX, and develop and check the displayed analytical arguments. Theorem, definition/assumption, remark and proof environments follow the AMS guide [11]. This declaration concerns preparation of the present text. It makes no attribution of completed work to an individual student, and is not a declaration that the team's solver has passed the course five-gate audit.

B.1 Analytical claim audits available in this draft

The following entries identify claims, risks, mathematical checks and their exact locations. They record the reasoning retained in this draft. No independent human sign-off or team verification record was supplied; the printed derivations and the cited external sources are the available evidence. These entries therefore should not be presented as completed student-performed verification records.

Claim 1: the ghost scheme has a second-order global upper bound. The assumptions are the smooth solution, fixed nonnegative β, exact Dirichlet values and the raw eliminated row in (3.10). The principal risk is identifying the order of the centred derivative with that of the eliminated PDE row. A falsifying algebraic observation would be a missing $2s_\beta/h$ term, an incorrect inward coefficient or a uniform consistency bound different from (4.6). The retained checks are exact exponential substitution in Corollary 4.2, the full-grid barrier in Proposition 4.6, and the separate energy route in Corollary 5.4. The decision is to retain an $O(h^2)$ theorem with its explicit assumptions. Its limit is that it asserts neither a measured slope nor a nonzero ghost leading-error constant.

Claim 2: a weighted matrix is eligible for CG. The assumptions are Dirichlet elimination, $N \geq 2$, $\beta \geq 0$ and the weights in (5.1). The risk is transferring continuous self-adjointness to the unweighted nonsymmetric matrix. A falsifier is a mismatched off-diagonal pair in $W_h^c A_h^c$ or a nonzero vector with nonpositive energy. The checks retained in Proposition 5.1 are coefficient pairing and an independently expanded edge-sum identity. The external interface requirement is the SciPy CG documentation [7]. The decision is to assert CG eligibility for K_h^c and positivity of its diagonal preconditioner. No execution, iteration count or efficiency result is asserted.

Claim 3: the FO exponential maximum error is genuinely first order. The assumptions are the exponential data and the exact FO discrete solution. The risk is interpreting an $O(h)$ upper bound as a sharp rate. A falsifier would be a failed sign in the interior defect, failure of the lower-boundary-defect inequality, or an incorrect comparison function boundary row. The evidence is the sequence of explicit inequalities in Theorem 4.8, ending in $C_\beta h \leq \max_j |e_{N,j}^{\mathrm{FO}}|$, combined with Theorem 4.7. The decision is to retain $\Theta(h)$ for this benchmark's boundary and global maximum norms. It is not extended to arbitrary solutions, to the ghost scheme or to a measured discrete L^2 slope.

B.2 Supplied assignment sources

The model, prescribed stencils and division of responsibilities come from the supplied course documents, rather than from an invented published application study. Their identification is retained here because no public URLs for these handouts were supplied:

- *problem_pack (3).pdf*, Direction 2, pp. 3–5: the square-domain mixed-boundary benchmark and the two required FDM closures.

- *Project_2_Brief (1)(1).pdf*, pp. 1–4: the project scope, report structure, evidence requirements, milestones and presentation allocation.
- *role_cards (2).pdf*, pp. 1–4: coordination and roles 1, 2 and 4, with roles 3 and 5 assigned separately.
- *Capacitor2d.pdf*, pp. 1–9: the separate microstrip FEM application, reviewed as contextual course material.
- *real_option_post.pdf*, pp. 1–8: the separate semilinear real-option/POST application, reviewed as contextual course material.
- *Numerical_PDEs_Assessment_Rubrics.pdf*, dated 11 September 2026, pp. 2–8: the previously supplied assessment descriptors.

The microstrip and real-option equations are not substituted for the selected Direction 2 model. The bibliography that follows contains only publicly identifiable sources actually cited in this manuscript, each with a clickable author, institutional, publisher or documentation link.